

Making USB-MIDI2 instruments with Zephyr

Beep-bloop ! I'm a kite

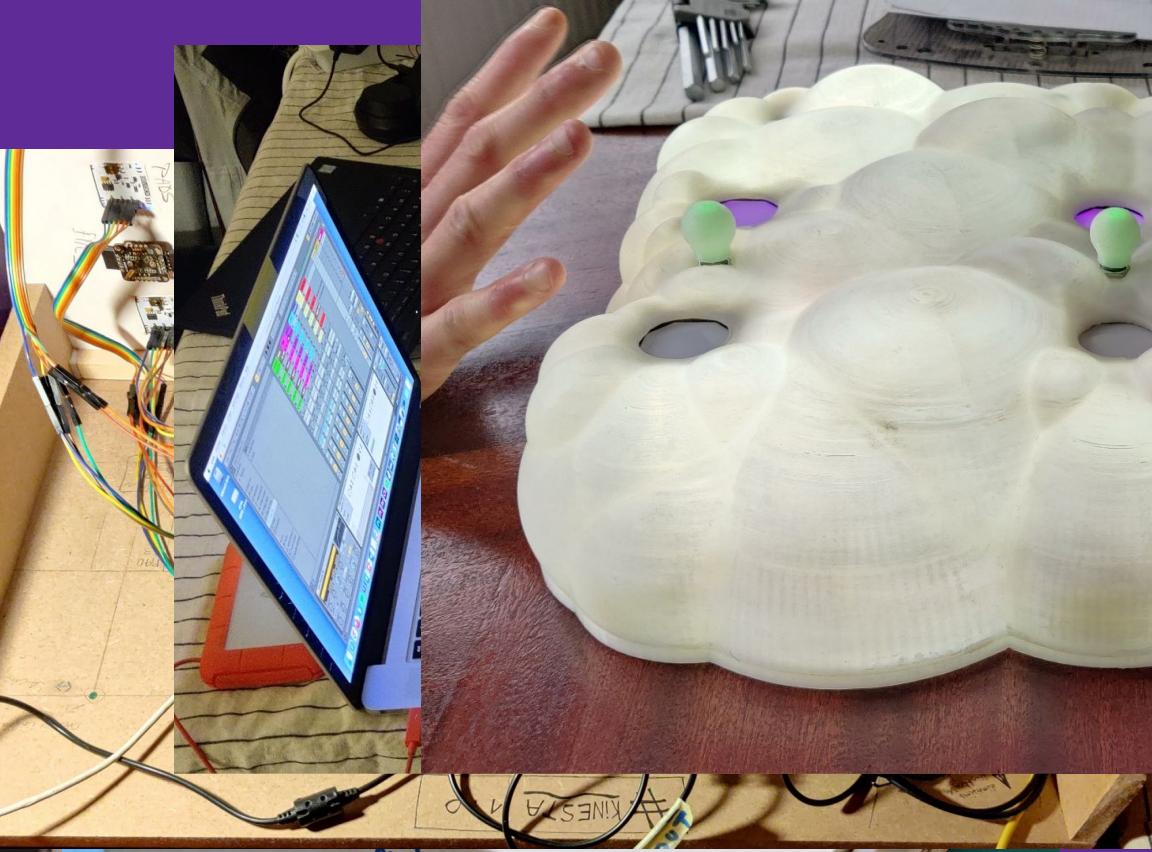
iTitou - Smartmonday March 2025



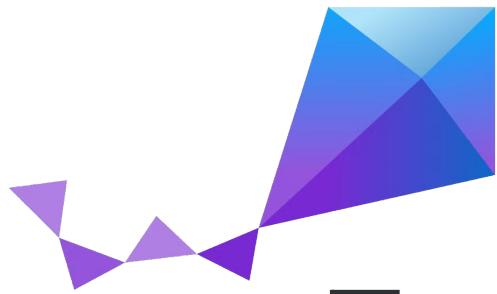
Who is he ?

- iTitou
 - Embedded & Backend developer
 - Ex-UrLab (~2013-2018) / FOSDEM delegate
 - Ex-ULB (BA1-BA3) / -VUB (MA1-MA2)
-
- Btw, I like trains





How it all started



Zephyr®

Enters the room

<https://zephyrproject.org>



[zephyrproject-rtos/zephyr](https://github.com/zephyrproject-rtos/zephyr)

Things running on Zephyr



What is Zephyr ?

- **Operating System for very small devices (microcontrollers)**
 - 100's MHz clock, "small" cores / 100's kB RAM
 - Dedicated peripherals (energy optimized)
- **Multi-platform**
- **Provides a framework for**
 - Multithreading with priorities
 - Networking (incl. BLE, Wifi, LoRa, etc...)
 - Security (TLS, crypto accelerators, etc...)
 - Device drivers
 - Utilities (logging, shell, etc...)
- **Higher level than "bare-metal" programming**
- **Compile-time configuration framework**
- **west command-line tool**

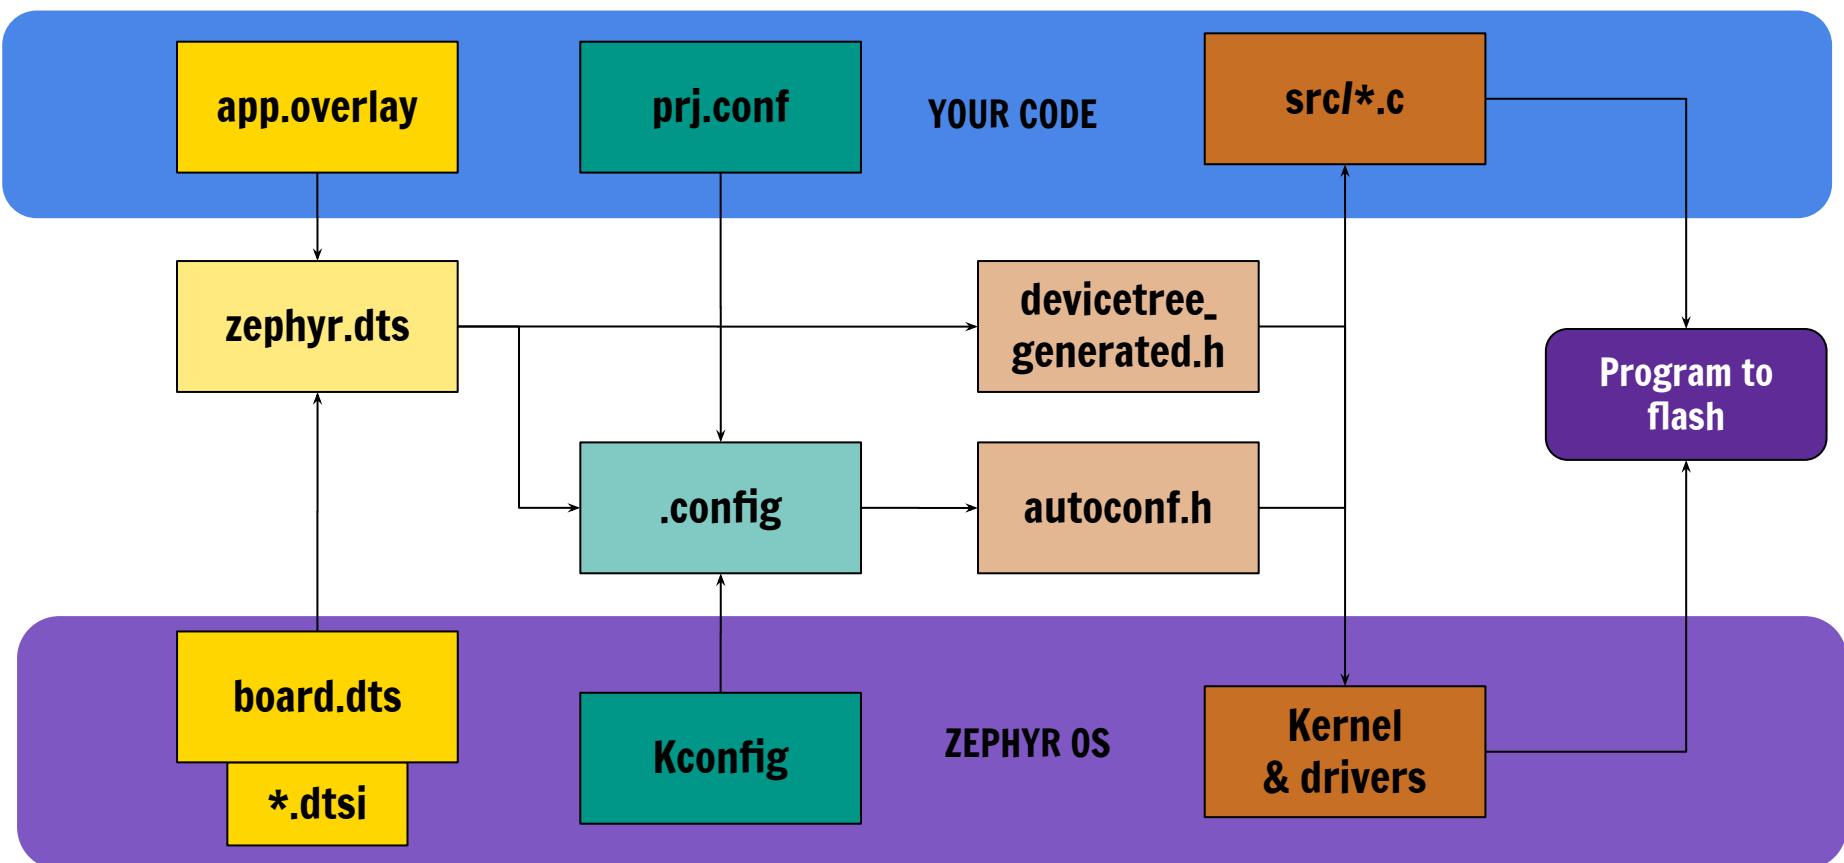


Zephyr vs. Arduino



	Zephyr	Arduino
Writing, compiling and running your code	Your IDE + west command line tool (integrate with VSCode; platformio)	Arduino IDE (platformio)
Adding external libraries	West modules	Copy / Paste libraries
Writing code for the target	Main function, threads, tasks and callback functions	setup() + loop() (callbacks with some libs)
Declaring devices	Device tree (KConfig for driver options)	C++ class instantiation in code
Accessing devices	Driver API with opaque const struct *device	C++ class accessors and methods

Building a Zephyr application



What is this Kconfig thing ?

- Originally, the (Linux) Kernel Config
- Define configurable items
 - Booleans (y/n)
 - Integers
 - Choices
 - Strings
- Available to Make
 - (what should be built ?)
- Available to C/C++ files (defines)

```
west build -t menuconfig
```

```
(Top) → Subsystems and OS Services → New USB device stack [EXPERIMENTAL]
Zephyr Kernel Configuration

Max compiled-in log level for usbd (Default) --->
(1024) USB device stack thread stack size (NEW)
(10) Maximum number of UDC events (NEW)
(8) Maximum number of USB device notification messages (NEW)
(1) USB device notification messages work delay (NEW)
[ ] USB Loopback Class (NEW)
[ ] Bluetooth HCI USB Transport Layer (NEW)
[ ] USB Mass Storage Device class support [EXPERIMENTAL] (NEW)
[ ] USB Audio 2 class support [EXPERIMENTAL] (NEW)
[*] USB MIDI 2.0 class support [EXPERIMENTAL]
    Max compiled-in log level for usbd midi2 (Default) --->
[ ] USB DFU Class support (NEW) ----

[Space/Enter] Toggle/enter  [ESC] Leave menu          [S] Save
[O] Load                  [?] Symbol info          [/] Jump to symbol
[F] Toggle show-help mode  [C] Toggle show-name mode [A] Toggle show-all mode
[Q] Quit (prompts for save) [D] Save minimal config (advanced)
```

What is this devicetree thing ?

- Originates from embedded Linux (again)
 - ! Compile-time in Zephyr vs. run-time in Linux
- Describe topology of all static peripherals
 - (ie. not dynamic like USB, PCI, ...)
- Includes & fragments
- Assemble, compile and decompile with dtc

=> When changing hardware: swap the device tree,
not the code !

```
/dts-v1;
#include <st/f4/stm32f407Xg.dtsci>
#include <st/f4/stm32f407(e-g)tx-pinctrl.dtsci>
#include <zephyr/dt-bindings/input/input-event-codes.h>

/ {
    model = "STMicroelectronics STM32F4DISCOVERY board";
    compatible = "st,stm32f4discovery";

    chosen {
        zephyr,console = &usart2;
        zephyr,shell-uart = &usart2;
        zephyr,sram = &sram0;
        zephyr,flash = &flash0;
        zephyr,ccm = &ccm0;
        zephyr,canbus = &can2;
    };

    leds {
        compatible = "gpio-leds";
        orange_led_3: led_3 {
            gpios = <&gpiod 13 GPIO_ACTIVE_HIGH>;
            label = "User LD3";
        };
        green_led_4: led_4 {
            gpios = <&gpiod 12 GPIO_ACTIVE_HIGH>;
            label = "User LD4";
        };
        red_led_5: led_5 {
            gpios = <&gpiod 14 GPIO_ACTIVE_HIGH>;
            label = "User LD5";
        };
        blue_led_6: led_6 {
            gpios = <&gpiod 15 GPIO_ACTIVE_HIGH>;
            label = "User LD6";
        };
    };

    gpio_keys {
        compatible = "gpio-keys";
        user_button: button {
            label = "Key";
            gpios = <&gpioa 0 GPIO_ACTIVE_HIGH>;
            zephyr,code = <INPUT_KEY_0>;
        };
    };
}
```

Creating a Zephyr application

- **Easiest way: find related example in samples/ and modify it to your needs**
- **Create from scratch:**
 - a. Create a new directory
 - b. Acc CMakeLists.txt (build rules)
 - c. Add prj.conf (Kconfig options)
 - d. (Optional) add app.overlay
 - e. Add C files in src/

zephyr / samples / subsys / usb /  

Add file ...

Name	Last commit message	Last commit date
..		
audio	usb: device: fix unaligned memory access in Audio c...	4 months ago
cdc_acm	samples: doc: drop obsolete modemanager notice	2 weeks ago
common	usb: device_next: fix Get Status request response	3 days ago
console-next	samples: usb: add console sample for the new USB ...	3 weeks ago
console	samples: usb: clean up legacy USB CDC ACM consol...	3 weeks ago
dfu-next	samples: usb: add sample for the new USB DFU imp...	last month
dfu	soc: mimxrt1180: Add USB support	last month
hid-keyboard	boards: mimxrt1060_evk: Convert to variants	3 months ago
hid-mouse	boards: nxp: add RT1060 EVKC support	2 months ago
inf	usb: cdc: add .inf file	6 years ago
mass	boards: mimxrt1060_evk: Convert to variants	3 months ago
midi	doc: usb: minor fixes to improve consistency	last month
shell	usb: device_next: add bulk transfers to loopback fun...	2 months ago
testusb	samples/subsys: Use hwmv2 native targets identifiers	last year
uac2_explicit_feedback	samples, tests: remove usage of space-separated lists	3 months ago
uac2_implicit_feedback	samples, tests: remove usage of space-separated lists	3 months ago
webusb-next	boards: mimxrt1060_evk: Convert to variants	3 months ago
webusb	samples: usb: fix API references	8 months ago
usb.rst	doc: samples: Adopt code-sample-category across tr...	6 months ago

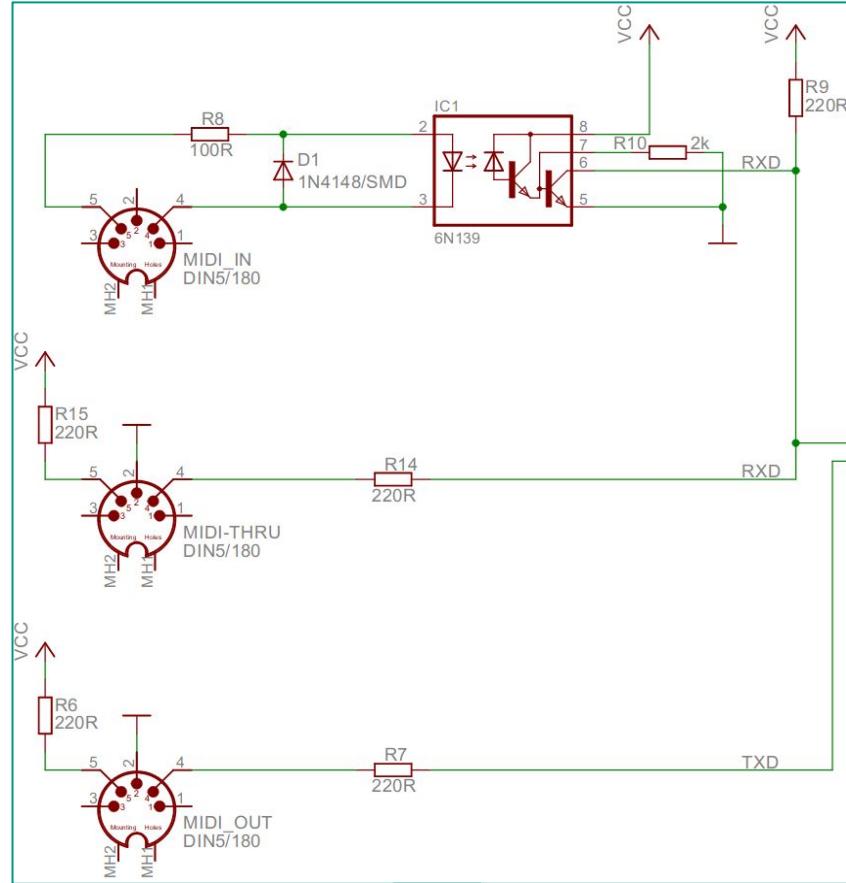
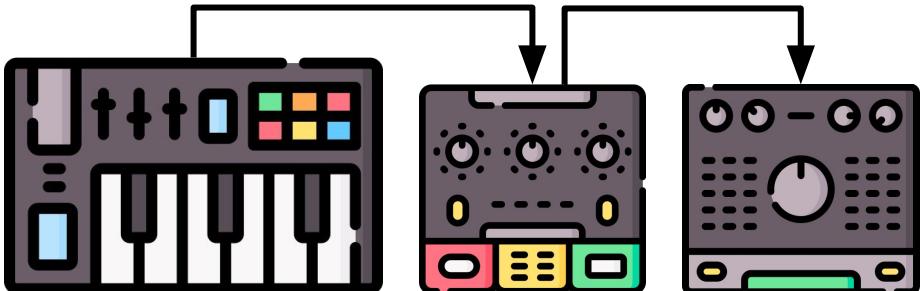
But wait...

There is no support for USB-MIDI in Zephyr

So let's build it !

MIDI (1) - physical layer

- “Musical scores” communication protocol
- Born in the 1980’s
- Iconic DIN5 connector
- Connect controllers to synthesizers
 - unidirectional + pass-through (daisy-chaining)
 - 16 channels
- UART @ 31.25 kb/s 8N1



MIDI (1) - messages

90 46 63 Play note A#4 on ch. 1

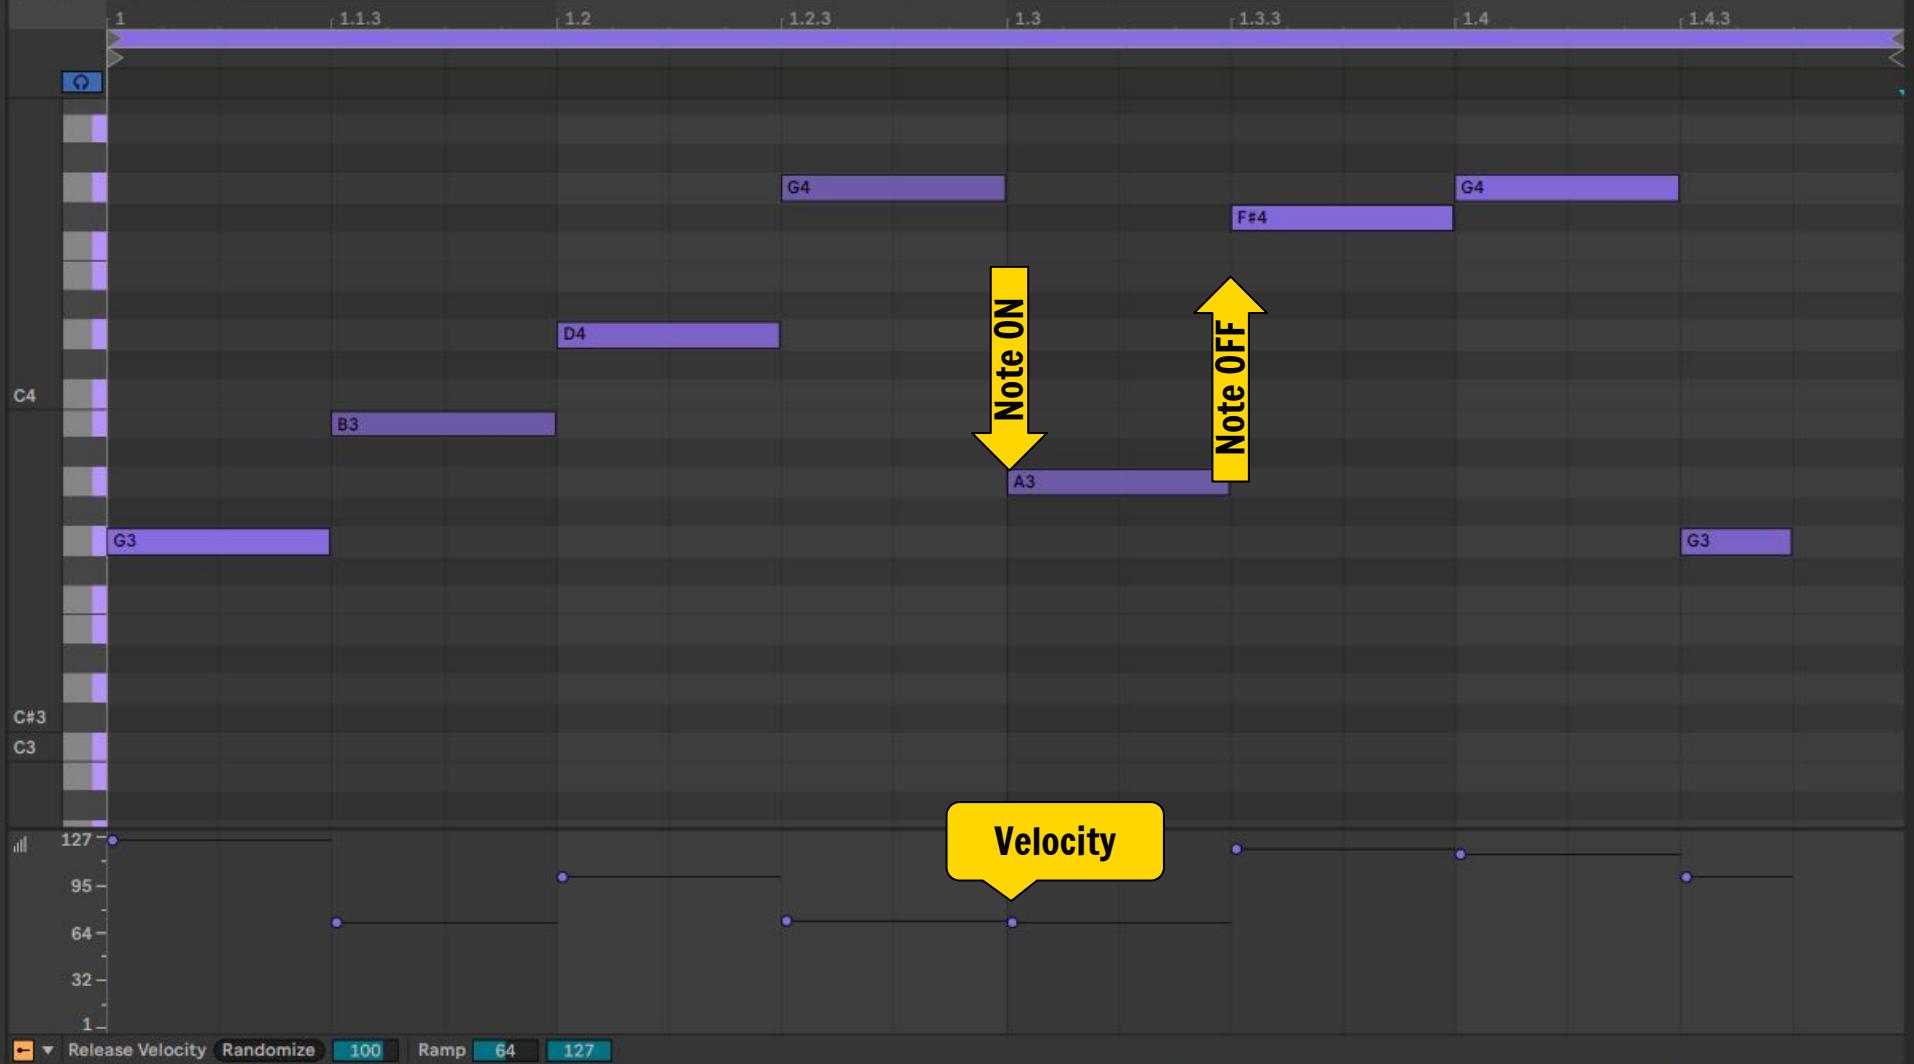
CF 2A Load program 42 on ch. 16

F0 46 4F 53 44 45 4D E7

Send ASCII “FOSDEM” as SysEx

- Variable length uint8 packet (1b start + 7b data)
 - Start of packet = msb is 1 otherwise always 0
- 16 chans, commands with 0, 1 or 2 params
- Var. length SysEx for proprietary payloads
- Send “music scores” events in real time
 - .mid file format with timing infos for replay

Some common MIDI 1 messages		Byte 0	Byte 1	Byte 2
Channel voice messages	Note off / on	<u>8/9</u>	Ch (0-F)	Note (0-7F)
	Control change	<u>B</u>	Ch (0-F)	Ctl number (0-7F)
	Program change	<u>C</u>	Ch (0-F)	Prog number (0-7F)
	Pitch bend (value = signed int14)	<u>E</u>	Ch (0-F)	Value LSB (0-7F)
System common	Start System Exclusive (SysEx)	<u>E</u>	0	(0-7F) [...] until “End of SysEx”
	End of System Exclusive	<u>E</u>	7	
System realtime	Clock (1 tick per quarter-note)	<u>E</u>	8	
	Start / Continue / Stop	<u>E</u>	A/B/C	



USB-MIDI 1.0 (1999)

- **MIDI jacks**
 - Represent USB-MIDI inputs / outputs
 - Embedded in/out = from/to host
 - External in/out = from/to the outside world
- **Elements for internal processing**
(synthesizers and effects)

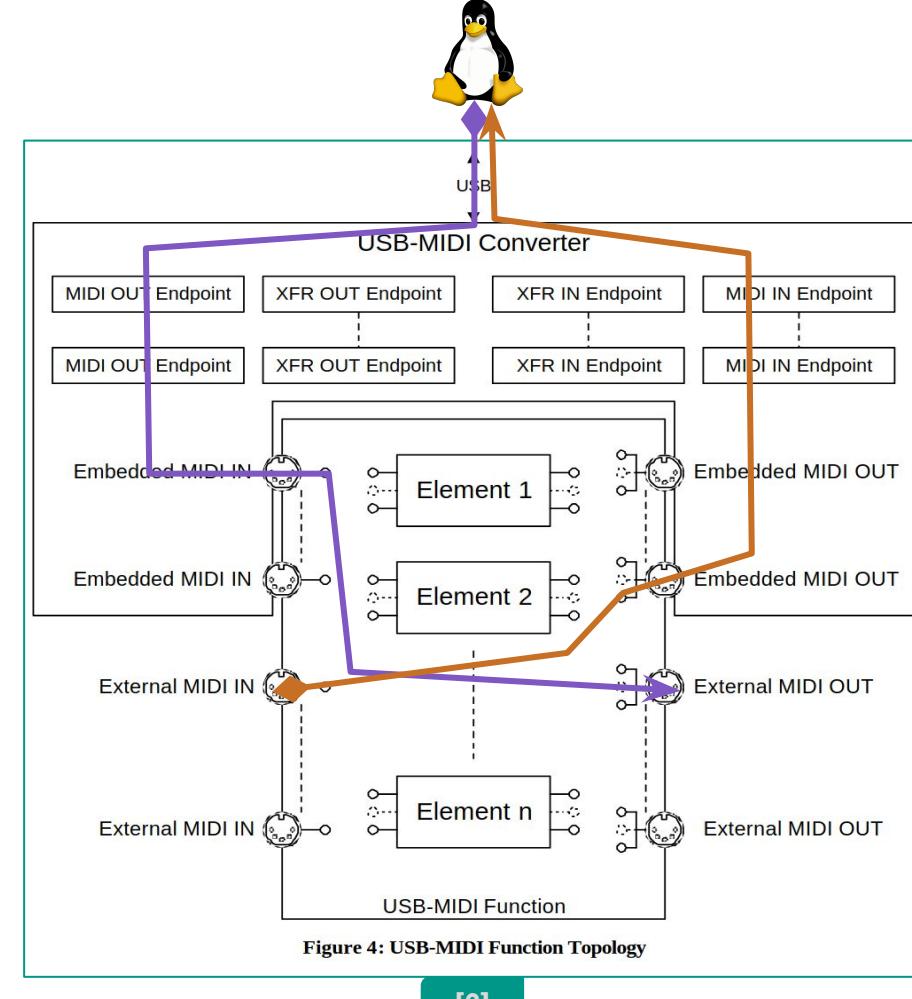
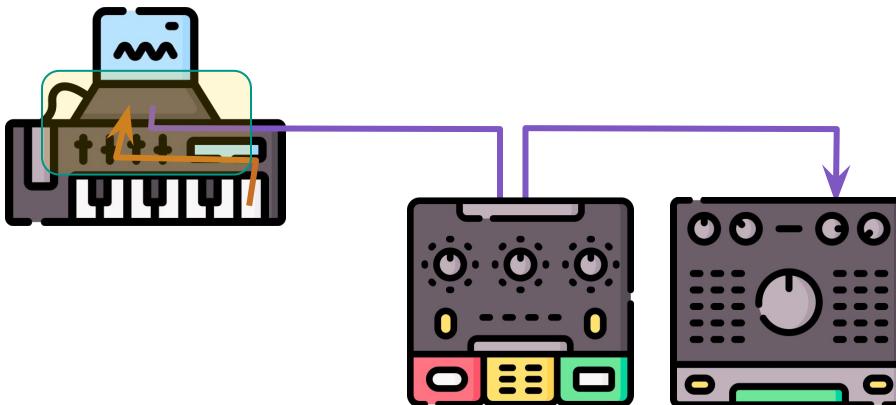
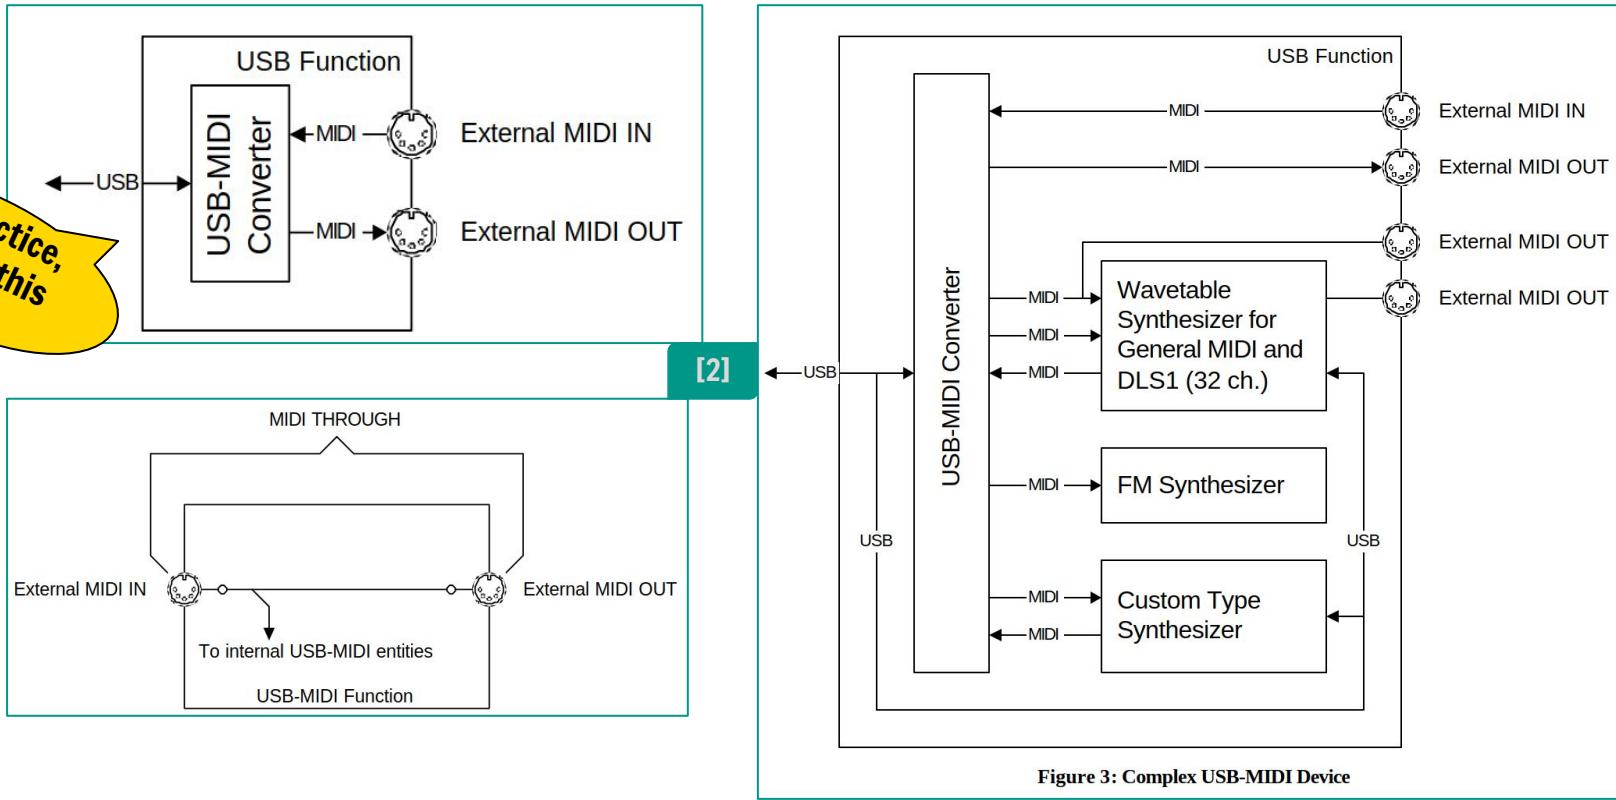


Figure 4: USB-MIDI Function Topology

USB-MIDI 1.0: Example topologies

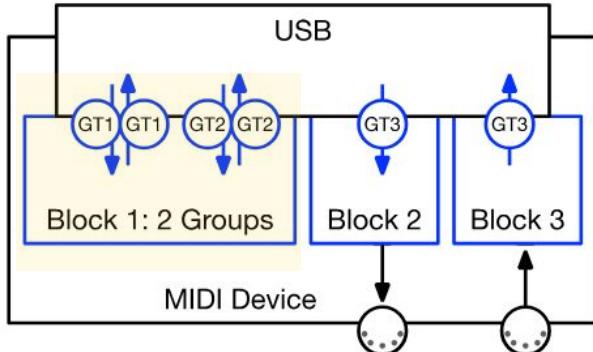


(*) At least in all consumer-grade devices I had in my hands

USB-MIDI 2.0 (2020)

- Simplified operational model
- Group Terminal Blocks
 - Functional relationship b/w ports

32 Channel MIDI Device plus external ports



It is not possible to guarantee MIDI-CI pairing with external MIDI jacks so IN and OUT jacks usually should not be members of a bidirectional Group Terminal Block.

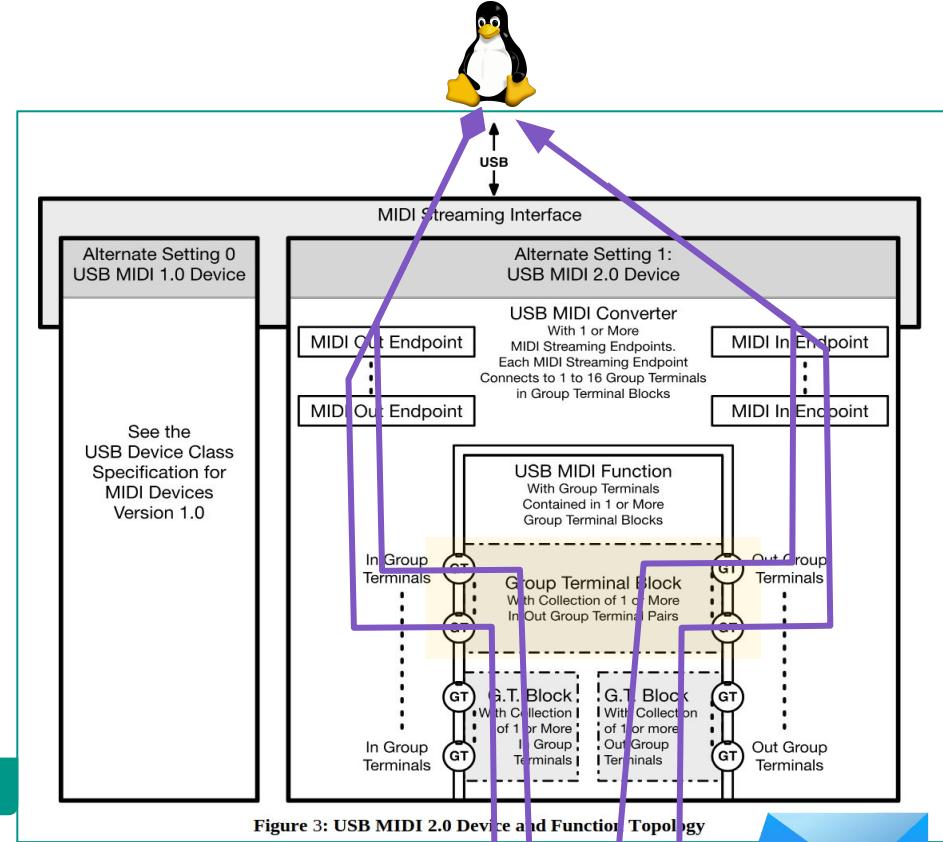
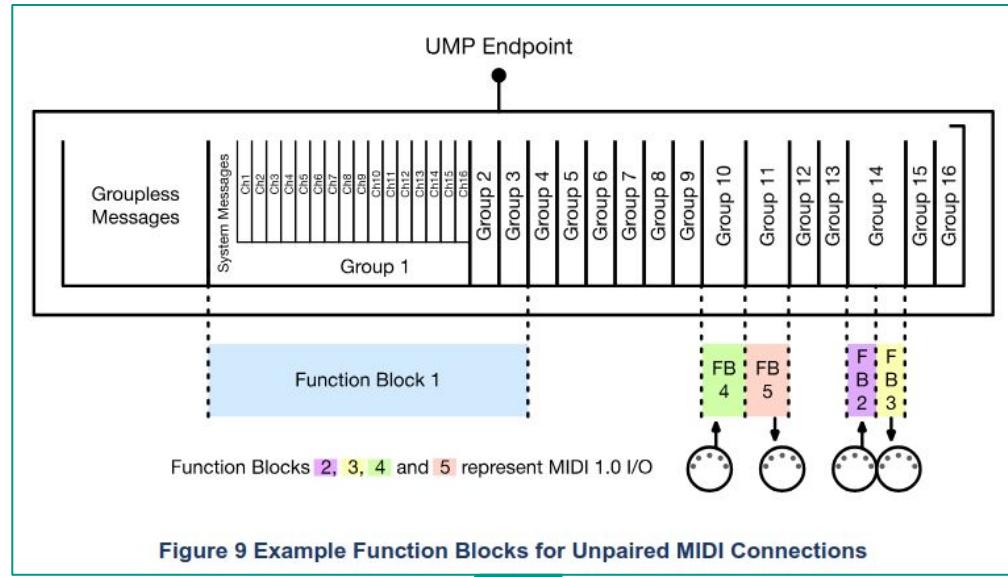


Figure 3: USB MIDI 2.0 Device and Function Topology

MIDI2 - Universal MIDI Packet endpoint

- New generic container format
 - 1 to 4 uint32
- Carry MIDI1 and MIDI2 messages:
 - Extended numerical ranges
 - Note attributes, tonal alterations
 - Tunings, chords, tempo,
 - time signature, lyrics, ...
- 16 Groups alloc. In Function Blocks
 - Discovery and dynamic configuration with UMP Stream (Groupless)



[4]

=> Acknowledge that most modern MIDI devices communicate on bidirectional, hi speed link (USB)

MIDI2

→ **Presentation**

is not

USB-MIDI2.0

→ **Transport**

is not

USB 2.0

→ **Levels below**

ex: send MIDI1 messages with USB-MIDI2 on USB1

Example Universal MIDI Packets



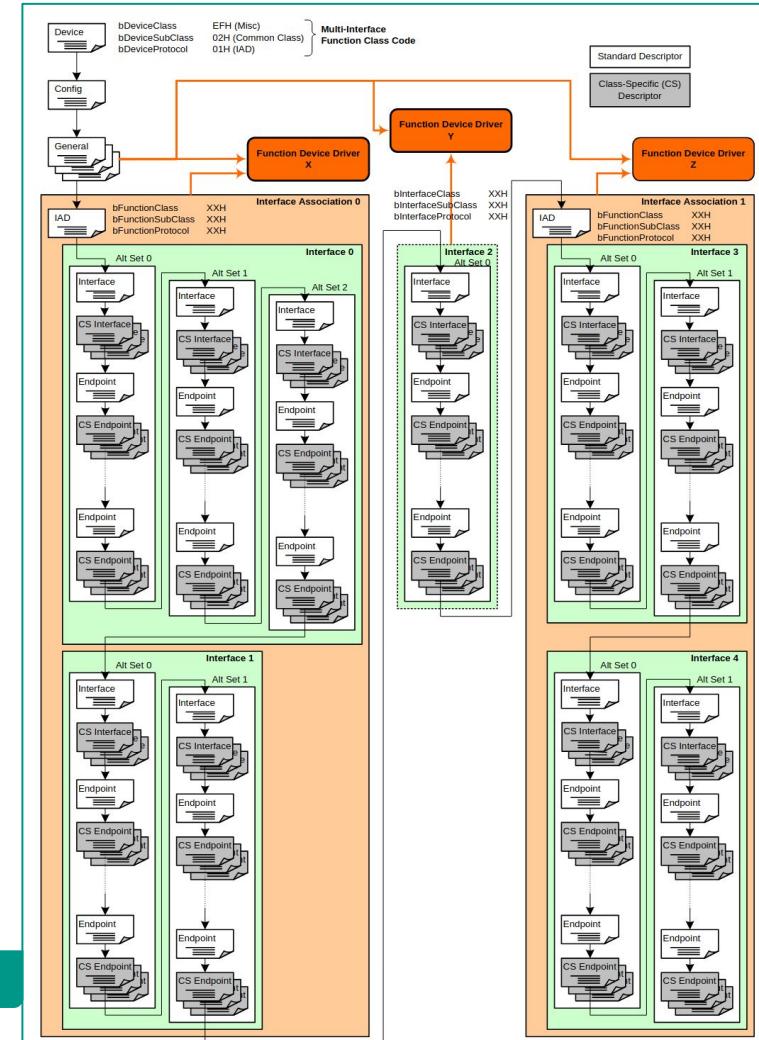
MIDI1 - Note Off / On (1 word - 32b)					
MT=2	Group (0-F)	8 / 9	Channel (0-F)	Note number (0-7F)	Velocity (0-7F)
	MSB				LSB

MIDI2 - Note Off / On (2 words - 64b)					
MT=4	Group (0-F)	8 / 9	Channel (0-F)	Note number (0-7F)	Attribute type
Velocity (0-FFFF)				Attribute value	
MSB					LSB

UMP Stream - Endpoint Name Notification (4 words - 128b)					
MT=F	Form (0-3)	Status = 0x03	Name [0]	Name [1]	Name [2..5]
					Name [6..9]
					Name [10..13]

USB descriptors, endpoints

- Descriptors
 - Expose capabilities to the host
 - Device class specific, iterable with bLength
 - “Key:value” mapping wrt. bDescriptor(Sub)Type
- Interfaces
 - Set of descriptors for one class “function”
- Endpoints
 - Unidirectional communication channels
 - Bulk, isochronous or interrupt
 - Declared in descriptors
- Class-specific requests
 - Out-of-band data exchange



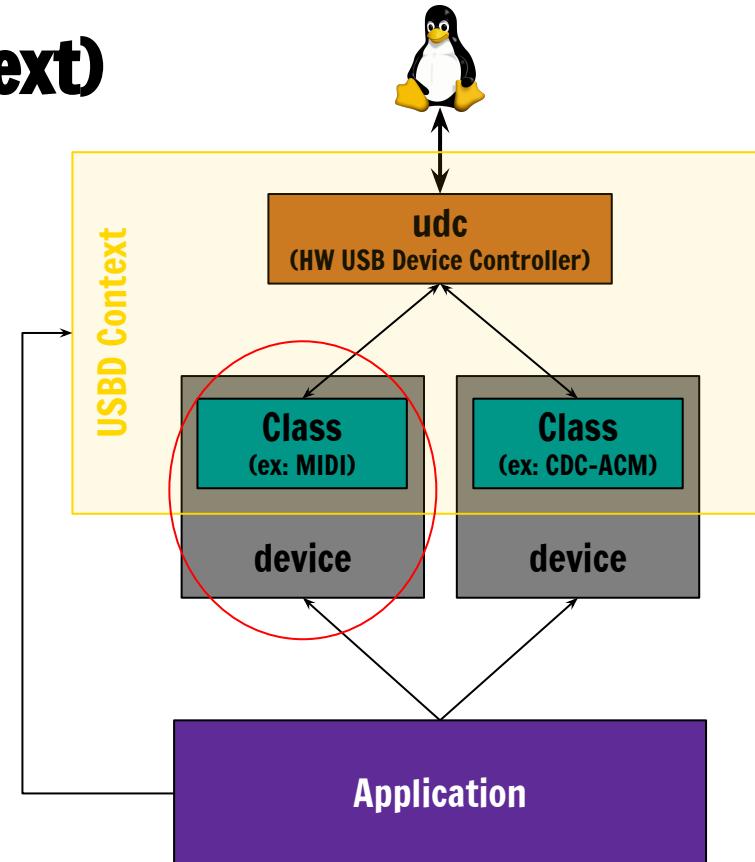
Our goal

Implement a USB-MIDI device

- Expose descriptors to the USB host (device capabilities)
 - USB class enumeration
- Respond to class-specific USB requests (topology)
 - Configured via the Zephyr device tree
- Exchange Universal MIDI packets over bulk USB endpoints
 - Runtime API

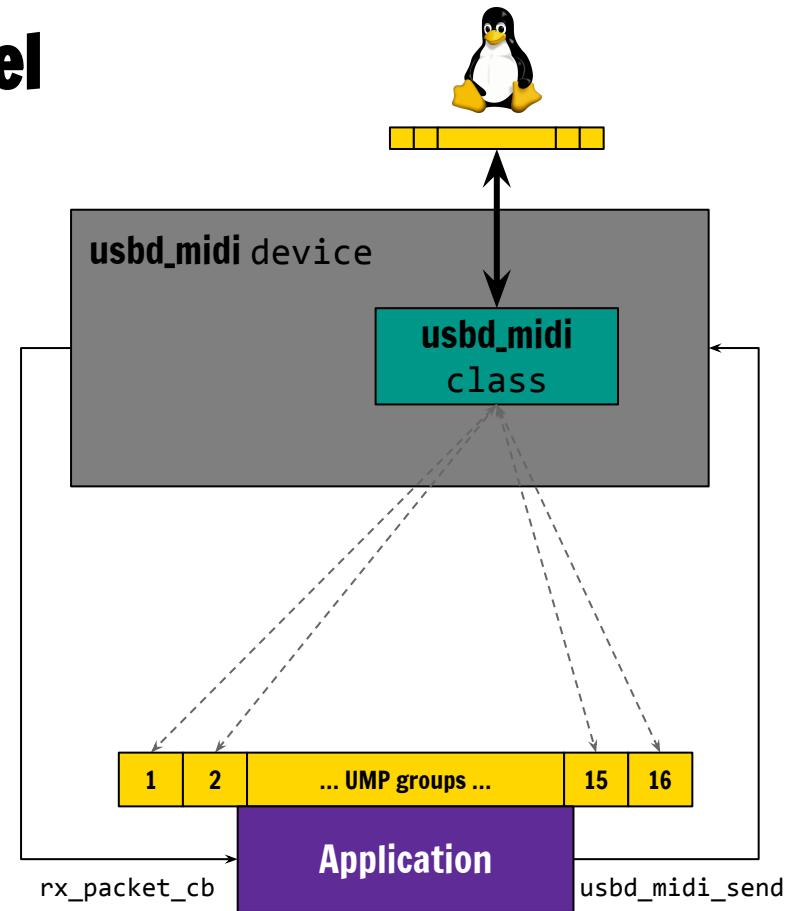
Zephyr udc, usbd (usb device_next)

- USBD Context manages wiring between
 - Low-level USB device controller
 - Registered USB device classes
 - Runtime state
- Easier than in old USB stack
 - Organize and pack descriptors
 - Handle multiple classes
- USB class drivers have to implement usbd_class_api
 - Corresponding Zephyr device



Zephyr usbd_midi: Operational model

- (up to) 16 UMP groups exchanged b/w host and Zephyr application
 - Group terminal blocks defined in device tree
- Send Universal MIDI Packet
 - Message Type (MT) field determines length
 - UMP Group field (except UMP_STREAM)
 - -> usbd_midi_send(dev, pkt)
- Receive Universal MIDI Packet
 - Register event handlers (all optional)
 - -> usbd_midi_set_ops(dev, {rx_packet_cb, ready_cb})
 - (avoid unneeded buffering)



Zephyr usbd_midi descriptors

- **Empty USB-MIDI1.0 on altsetting 0**
 - (ie. no jacks or elements)
 - Required by the spec, won't enumerate without
- **USB-MIDI2.0 on altsetting 1**
 - Group terminal blocks (topology) retrieved via separate class-specific request (as per spec.)
 - MUST be selected by host
- **One pair of endpoints per interface**

Class			Interface association
	If 0	Alt 0	Audiocontrol Interface
			Audio class-specific
			MIDIStreaming Interface 1.0
			MIDIStreaming class-specific
			Endpoint
			Endpoint class-specific
			Endpoint
			Endpoint class-specific
			MIDIStreaming Interface 2.0
			MIDIStreaming class-specific
			Endpoint
			Endpoint class-specific
			Endpoint
			Endpoint class-specific

**Let's make
some noise !**

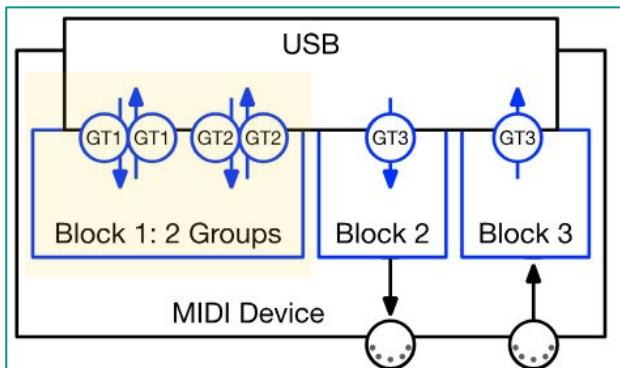


Set up usbd_midi in Zephyr application

app.overlay

prj.conf

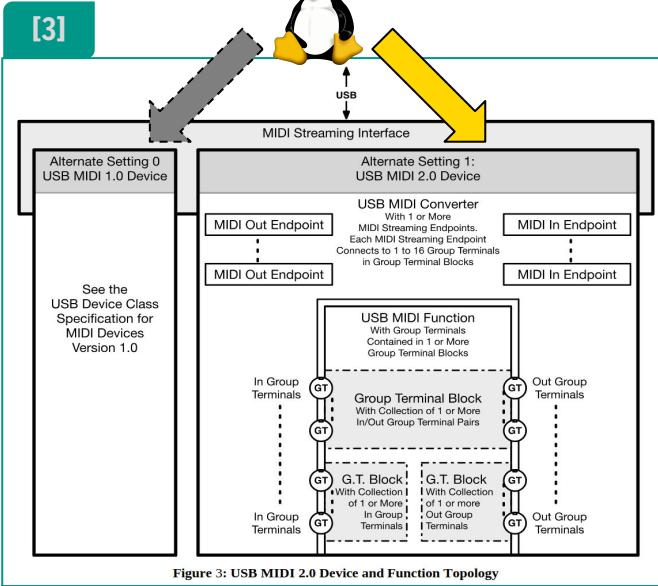
```
CONFIG_USB_DEVICE_STACK_NEXT=y  
CONFIG_USBD_MIDI2_CLASS=y
```



```
/ {  
    usb_midi: usb_midi {  
        compatible = "zephyr,midi2-device";  
        status = "okay";  
        #address-cells = <1>;  
        #size-cells = <1>;  
  
        block1 {  
            reg = <0 2>;  
            protocol = "midi2";  
            terminal-type = "bidirectional";  
            status = "okay";  
        };  
        /* ... */  
        block3 {  
            reg = <2 1>;  
            protocol = "midi1-up-to-64b";  
            terminal-type = "input-only";  
            status = "okay";  
        };  
    };
```

Register application event handlers

- Register callbacks
 - MIDI packet from host
 - Connection status change



```
#include <zephyr/drivers/gpio.h>
#include <zephyr/usb/class/usbd_midi2.h>

static void on_device_ready(const struct device *dev,
                           const bool ready)
{
    gpio_pin_set_dt(&led, ready);
}

static const struct usbd_midi_ops ops = {
    .ready_cb = on_device_ready,
    .rx_packet_cb = rx_midi,
};

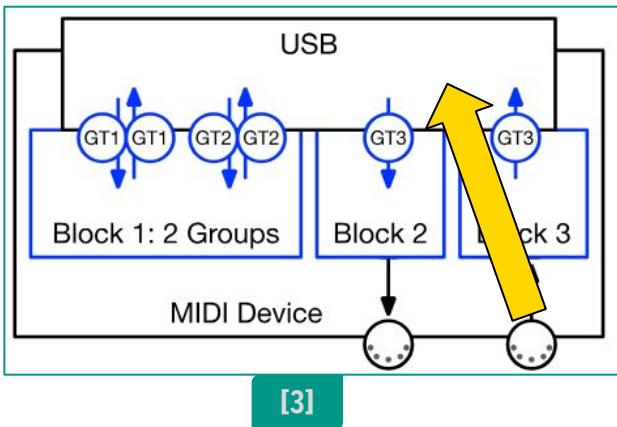
int main(void)
{
    /* ... Initialize application state ... */

    usbd_midi_set_ops(midi, &ops);
}
```

More in 2 slides...

MIDI1 keyboard

- Leverage input subsys
- Build UMP, then send it



```
#include <zephyr/input/input.h>
#include <zephyr/usb/class/usbd_midi2.h>

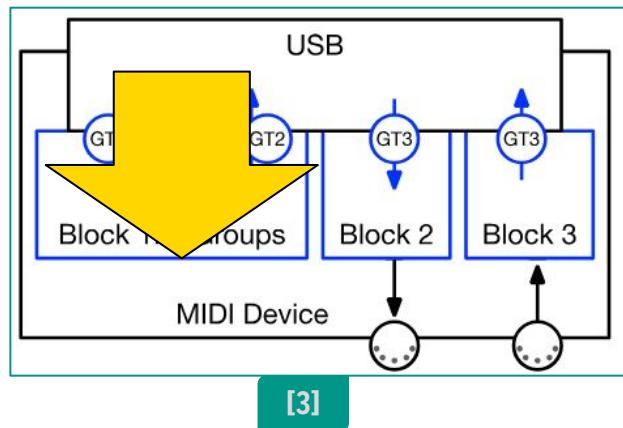
static const struct device *usb_midi =
    DEVICE_DT_GET(DT_NODELABEL(usb_midi));

static void key_press(struct input_event *evt, void *_)
{
    if (evt->type != INPUT_EV_KEY || evt->code > 0x7f) {
        return;
    }
    uint8_t command = evt->value
                    ? UMP_MIDI_NOTE_ON
                    : UMP_MIDI_NOTE_OFF; Key press
    uint8_t note = evt->code;
    const struct midi_ump ump = UMP_MIDI1_CHANNEL_VOICE(
        UMP_group 3, UMP group 3 command, 0, MIDI chan 1
        note, 100 Velocity ≈ “fortissimo”
    );
    usbd_midi_send(usb_midi, ump);
}

INPUT_CALLBACK_DEFINE(NULL, key_press, NULL);
```

MIDI2 Synthesizer

- React to MIDI2 Note On/Off sent by the host



```
static void rx_midi(const struct device *dev,
                    const struct midi_ump pkt)
{
    if (UMP_MT(pkt) != UMP_MT_MIDI2_CHANNEL_VOICE) {
        return;
    }

    uint8_t channel = (UMP_GROUP(pkt) << 4)
                    | UMP_MIDI_CHANNEL(pkt);
    if (channel >= 32)
        return;

    switch (UMP_MIDI_COMMAND(pkt)) {
        case UMP_MIDI_NOTE_ON:
            start_tone(channel,
                        UMP_MIDI1_P1(pkt),
                        UMP_MIDI2_VELOCITY(pkt));
            break;
        case UMP_MIDI_NOTE_OFF:
            stop_tone(channel);
            break;
    }
}
```

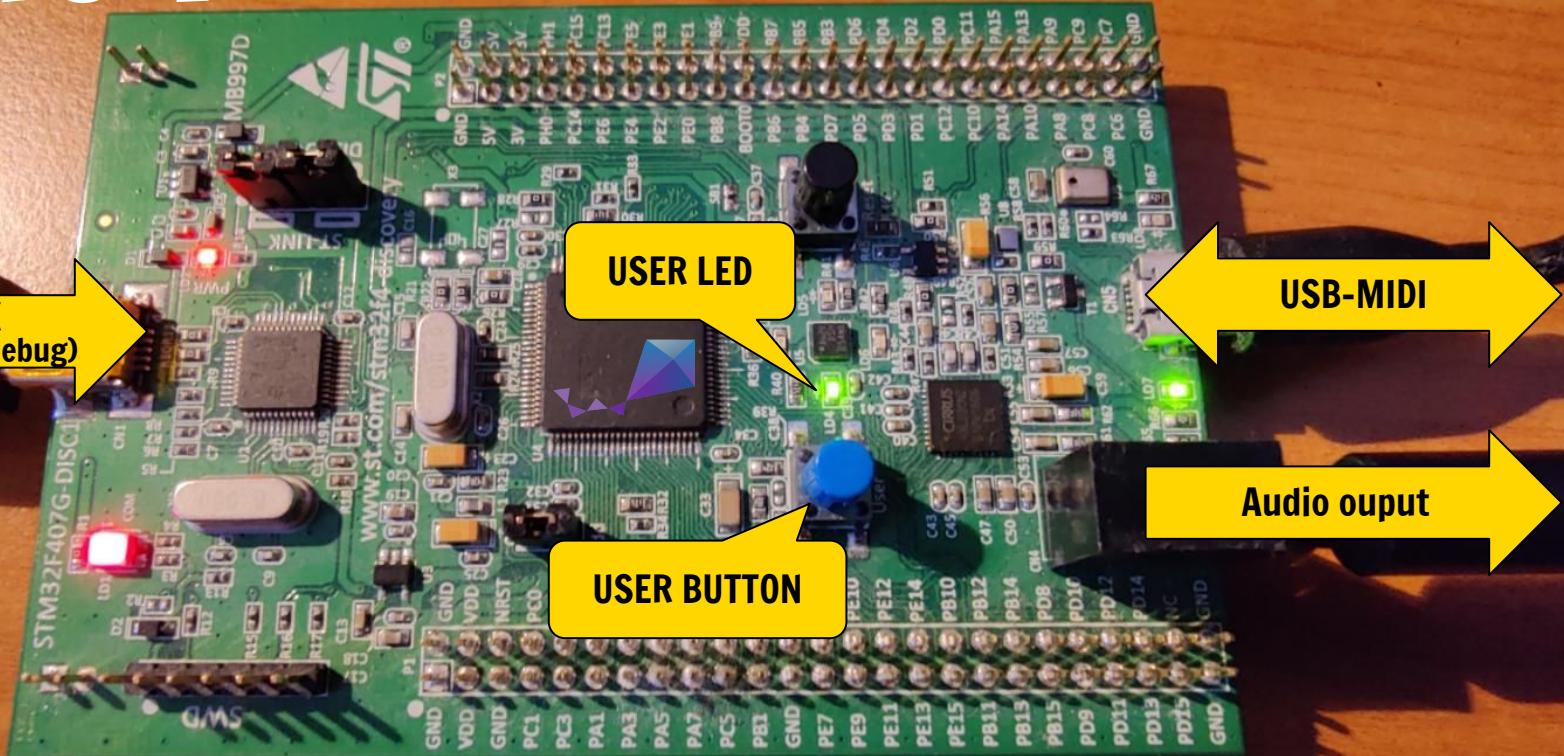
2 groups
x16 chans

Note pitch

Note volume

Let's just assume we have these

Demo!





USB- MIDI2 is fairly new

Let's build fancy instruments and controllers in the open !

(Coming in Zephyr 4.1)

Thank you !

Going further and beyond

- [Pull Request “add new MIDI 2.0 device class” in Zephyr](#)
- [MIDI 2.0 on Linux \(kernel documentation\)](#)
- [USB Device Support in Zephyr RTOS from the Application Perspective - Johann Fischer \(Video from EOSS 2023\)](#)
- [Open Source MIDI driver in Windows](#)

- [World's First MIDI Shellcode :: portasynthinc3's blog](#)

Figures and images

[1] [Olimex SHIELD MIDI schematics](#)

[2] [USB Device Class Definition for MIDI Devices](#)

[3] [USB Device Class Definition for MIDI Devices Release 2](#)

[4] [Universal MIDI Packet \(UMP\) Format and MIDI 2.0 Protocol](#)

[5] [USB Interface Association Descriptor Device Class Code and Use Model](#)

[*] [Flaticon “Audio Edition” pack](#)

[*] https://commons.wikimedia.org/wiki/File:Orange_Banjo.svg